

through better data filtering. Instead, it is a fundamental question of geometric control in high-dimensional curved spaces, with profound implications for open-weight model deployment. The structural fragility we expose cannot be eliminated through pre-deployment alignment alone, requiring the community to reconsider whether current deployment paradigms are compatible with safety requirements.

References

- [1] ALKAEED, M., ABIOYE, S., QAYYUM, A., MEKKI, Y. M., BERROU, I., ABDALLAH, M., AL-FUQAHA, A., BILAL, M., AND QADIR, J. Open foundation models in healthcare: Challenges, paradoxes, and opportunities with genai driven personalized prescription. *arXiv preprint arXiv:2502.04356* (2025).
- [2] AMARI, S.-I. Information geometry of neural networks—an overview—. *Mathematics of Neural Networks: Models, Algorithms and Applications* (1997), 15–23.
- [3] ARDITI, A., OBESO, O., SYED, A., PALEKA, D., PANICKSSERY, N., GURNEE, W., AND NANDA, N. Refusal in language models is mediated by a single direction. *Advances in Neural Information Processing Systems 37* (2024), 136037–136083.
- [4] ARORA, S., AND GOYAL, A. A theory for emergence of complex skills in language models. *arXiv preprint arXiv:2307.15936* (2023).
- [5] ARTURI, D. A. R., ZHANG, E., ANSAH, A. A., ZHU, K., PANDA, A., AND BALWANI, A. Shared parameter subspaces and cross-task linearity in emergently misaligned behaviour. In *Mechanistic Interpretability Workshop at NeurIPS 2025* (2025).
- [6] BETLEY, J., WARNCKE, N., SZTYBER-BETLEY, A., TAN, D., BAO, X., SOTO, M., SRIVASTAVA, M., LABENZ, N., AND EVANS, O. Training large language models on narrow tasks can lead to broad misalignment. *Nature* 649, 8097 (2026), 584–589.
- [7] BHANDARI, K. R., CHEN, P.-Y., AND GAO, J. Forecasting open-weight AI model growth on huggingface. *arXiv preprint arXiv:2502.15987* (2025).
- [8] BRICKEN, T., TEMPLETON, A., BATSON, J., CHEN, B., JERMYN, A., CONERLY, T., TURNER, N., ANIL, C., DENISON, C., ASKELL, A., LASENBY, R., WU, Y., KRAVEC, S., SCHIEFER, N., MAXWELL, T., JOSEPH, N., HATFIELD-DODDS, Z., TAMKIN, A., NGUYEN, K., MCLEAN, B., BURKE, J. E., HUME, T., CARTER, S., HENIGHAN, T., AND OLAH, C. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread* (2023). <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- [9] BURTON, J. W., LOPEZ-LOPEZ, E., HECHTLINGER, S., RAHWAN, Z., AESCHBACH, S., BAKKER, M. A., BECKER, J. A., BERDITCHEVSKAIA, A., BERGER, J., BRINKMANN, L., ET AL. How large language models can reshape collective intelligence. *Nature human behaviour* 8, 9 (2024), 1643–1655.
- [10] CHOKSI, M. Z., MANDEL, I., AND BENTHALL, S. The brief and wondrous life of open models. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency* (2025), pp. 3224–3240.
- [11] CHUA, K., LEI, Q., AND LEE, J. D. How fine-tuning allows for effective meta-learning. *Advances in Neural Information Processing Systems 34* (2021), 8871–8884.
- [12] COBBE, K., KOSARAJU, V., BAVARIAN, M., CHEN, M., JUN, H., KAISER, L., PLAPPERT, M., TWOREK, J., HILTON, J., NAKANO, R., ET AL. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168* (2021).
- [13] COMANICI, G., BIEBER, E., SCHAEKERMANN, M., PASUPAT, I., SACHDEVA, N., DHILLON, I., BLISTEIN, M., RAM, O., ZHANG, D., ROSEN, E., ET AL. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261* (2025).
- [14] COVER, T. M. *Elements of information theory*. John Wiley & Sons, 1999.
- [15] CUNNINGHAM, H., EWART, A., RIGGS, L., HUBEN, R., AND SHARKEY, L. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600* (2023).

- [16] DAS, A., BORAH, A., JAIN, V., AND CHADHA, A. Alignguard-lora: Alignment-preserving fine-tuning via fisher-guided decomposition and riemannian-geodesic collision regularization. *arXiv preprint arXiv:2508.02079* (2025).
- [17] DAVIS, C., AND KAHAN, W. M. The rotation of eigenvectors by a perturbation. iii. *SIAM Journal on Numerical Analysis* 7, 1 (1970), 1–46.
- [18] DONG, Y., MU, R., ZHANG, Y., SUN, S., ZHANG, T., WU, C., JIN, G., QI, Y., HU, J., MENG, J., ET AL. Safeguarding large language models: A survey. *Artificial Intelligence Review* 58, 12 (2025), 382.
- [19] DU, S. S., HU, W., KAKADE, S. M., LEE, J. D., AND LEI, Q. Few-shot learning via learning the representation, provably. *arXiv preprint arXiv:2002.09434* (2020).
- [20] FRASER, K. C., DAWKINS, H., NEJADGHOLO, I., AND KIRITCHENKO, S. Fine-tuning lowers safety and disrupts evaluation consistency. *arXiv preprint arXiv:2506.17209* (2025).
- [21] GAO, L., LA TOUR, T. D., TILLMAN, H., GOH, G., TROLL, R., RADFORD, A., SUTSKEVER, I., LEIKE, J., AND WU, J. Scaling and evaluating sparse autoencoders. In *The Thirteenth International Conference on Learning Representations* (2025).
- [22] GLIWA, B., MOCHOL, I., BIESEK, M., AND WAWER, A. SAMSum corpus: A human-annotated dialogue dataset for abstractive summarization. In *Proceedings of the 2nd Workshop on New Frontiers in Summarization* (Hong Kong, China, Nov. 2019), Association for Computational Linguistics, pp. 70–79.
- [23] GRATTAFFIORI, A., DUBEY, A., JAUHRI, A., PANDEY, A., KADIAN, A., AL-DAHLE, A., LETMAN, A., MATHUR, A., SCHELLEN, A., VAUGHAN, A., ET AL. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783* (2024).
- [24] GUAN, Z., HU, M., ZHU, R., LI, S., AND VULLIKANTI, A. Benign samples matter! fine-tuning on outlier benign samples severely breaks safety. In *Forty-second International Conference on Machine Learning* (2025).
- [25] HE, L., XIA, M., AND HENDERSON, P. What is in your safe data? Identifying benign data that breaks safety. In *First Conference on Language Modeling* (2024).
- [26] HSIUNG, L., PANG, T., TANG, Y.-C., SONG, L., HO, T.-Y., CHEN, P.-Y., AND YANG, Y. Your task may vary: A systematic understanding of alignment and safety degradation when fine-tuning LLMs.
- [27] HU, E. J., SHEN, Y., WALLIS, P., ALLEN-ZHU, Z., LI, Y., WANG, S., WANG, L., CHEN, W., ET AL. Lora: Low-rank adaptation of large language models. *ICLR* 1, 2 (2022), 3.
- [28] JI, J., LIU, M., DAI, J., PAN, X., ZHANG, C., BIAN, C., CHEN, B., SUN, R., WANG, Y., AND YANG, Y. Beavertails: Towards improved safety alignment of LLM via a human-preference dataset. *Advances in Neural Information Processing Systems* 36 (2023), 24678–24704.
- [29] JIANG, X., WANG, X., AND STICH, S. U. Loram: Low-rank adaptation of large language models on manifold. In *Sparsity in LLMs (SLLM): Deep Dive into Mixture of Experts, Quantization, Hardware, and Inference*.
- [30] JOHNSON, W. B., LINDENSTRAUSS, J., ET AL. Extensions of lipschitz mappings into a hilbert space. *Contemporary mathematics* 26, 189–206 (1984), 1.
- [31] KAPOOR, S., BOMMASANI, R., KLYMAN, K., LONGPRE, S., RAMASWAMI, A., CIHON, P., HOPKINS, A., BANKSTON, K., BIDERMAN, S., BOGEN, M., ET AL. On the societal impact of open foundation models. *arXiv preprint arXiv:2403.07918* (2024).
- [32] KAUSHIK, P., CHAUDHARI, S., VAIDYA, A., CHELLAPPA, R., AND YUILLE, A. The universal weight subspace hypothesis, 2025.
- [33] KUNSTNER, F., HENNIG, P., AND BALLES, L. Limitations of the empirical fisher approximation for natural gradient descent. *Advances in neural information processing systems* 32 (2019).
- [34] LEE, A., BAI, X., PRES, I., WATTENBERG, M., KUMMERFELD, J. K., AND MIHALCEA, R. A mechanistic understanding of alignment algorithms: A case study on DPO and toxicity. In *Forty-first International Conference on Machine Learning* (2024).

- [35] LI, C., FARKHOOR, H., LIU, R., AND YOSINSKI, J. Measuring the intrinsic dimension of objective landscapes. *arXiv preprint arXiv:1804.08838* (2018).
- [36] LIU, Z., NIE, Y., TAN, Y., YUE, X., CUI, Q., WANG, C., ZHU, X., AND ZHENG, B. Safety alignment for vision language models. *arXiv preprint arXiv:2405.13581* (2024).
- [37] MALLADI, S., WETTIG, A., YU, D., CHEN, D., AND ARORA, S. A kernel-based view of language model fine-tuning. In *International Conference on Machine Learning* (2023), PMLR, pp. 23610–23641.
- [38] MANTRI, K. S. I., ELIASOF, M., SCHÖNLIEB, C.-B., AND RIBEIRO, B. Rethinking fine-tuning through geometric perspective. In *UniReps: 2nd Edition of the Workshop on Unifying Representations in Neural Models* (2024).
- [39] NIU, Y., XIAO, H., LIU, D., CHEN, N., AND LI, J. Mitigating the safety alignment tax with null-space constrained policy optimization. *arXiv preprint arXiv:2512.11391* (2025).
- [40] OSBORNE, C., DING, J., AND KIRK, H. R. The ai community building the future? a quantitative analysis of development activity on hugging face hub. *Journal of Computational Social Science* 7, 2 (2024), 2067–2105.
- [41] PARK, S. M., GEORGIEV, K., ILYAS, A., LECLERC, G., AND MKADRY, A. Trak: attributing model behavior at scale. In *Proceedings of the 40th International Conference on Machine Learning* (2023), ICML’23, JMLR.org.
- [42] PENG, S. Y., CHEN, P.-Y., HULL, M., AND CHAU, D. H. Navigating the safety landscape: Measuring risks in finetuning large language models. *Advances in Neural Information Processing Systems* 37 (2024), 95692–95715.
- [43] QI, X., WEI, B., CARLINI, N., HUANG, Y., XIE, T., HE, L., JAGIELSKI, M., NASR, M., MITTAL, P., AND HENDERSON, P. On evaluating the durability of safeguards for open-weight LLMs. *arXiv preprint arXiv:2412.07097* (2024).
- [44] QI, X., ZENG, Y., XIE, T., CHEN, P.-Y., JIA, R., MITTAL, P., AND HENDERSON, P. Fine-tuning aligned language models compromises safety, even when users do not intend to! *arXiv preprint arXiv:2310.03693* (2023).
- [45] RAJABI, S., AND RAMBHATLA, S. Enhancing fine-tuning efficiency of LLMs through gradient subspace tracking. In *Adaptive Foundation Models: Evolving AI for Personalized and Efficient Learning* (2024).
- [46] RIFAI, S., DAUPHIN, Y. N., VINCENT, P., BENGIO, Y., AND MULLER, X. The manifold tangent classifier. *Advances in neural information processing systems* 24 (2011).
- [47] SAGUN, L., EVCI, U., GUNAY, V. U., DAUPHIN, Y., AND BOTTOU, L. Empirical analysis of the hessian of over-parametrized neural networks. *arXiv preprint arXiv:1706.04454* (2017).
- [48] SCHULMAN, J., LEVINE, S., ABBEEL, P., JORDAN, M., AND MORITZ, P. Trust region policy optimization. In *International conference on machine learning* (2015), PMLR, pp. 1889–1897.
- [49] SCHULMAN, J., WOLSKI, F., DHARIWAL, P., RADFORD, A., AND KLIMOV, O. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347* (2017).
- [50] SHUTTLEWORTH, R. S., ANDREAS, J., TORRALBA, A., AND SHARMA, P. LoRA vs full fine-tuning: An illusion of equivalence. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems* (2025).
- [51] SOLIGO, A., TURNER, E., RAJAMANOCHARAN, S., AND NANDA, N. Convergent linear representations of emergent misalignment. *arXiv preprint arXiv:2506.11618* (2025).
- [52] TAORI, R., GULRAJANI, I., ZHANG, T., DUBOIS, Y., LI, X., GUESTRIN, C., LIANG, P., AND HASHIMOTO, T. B. Stanford alpaca: An instruction-following llama model. https://github.com/tatsu-lab/stanford_alpaca, 2023.
- [53] TEAM, Q. Qwen3 technical report, 2025.
- [54] THIRUTHUMMAL, A. A., KIM, E.-J., AND SHELKAG, S. Information geometry of evolution of neural network parameters while training. *Neurocomputing* 597 (2024), 128007.
- [55] TOUVRON, H., LAVRIL, T., IZACARD, G., MARTINET, X., LACHAUX, M.-A., LACROIX, T., ROZIÈRE, B., GOYAL, N., HAM-BRO, E., AZHAR, F., ET AL. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971* (2023).